

BRAIN PLUG

Multi-Tenant AI Agent SaaS Platform & Real-Time RAG Streaming Engine

A production-grade cloud architecture detailing dynamic tenant scoping, dense vector knowledge ingestion, Shadow DOM chat widget encapsulation, AES-256-GCM cryptography, and Customer Relations (CR) ticketing workflows.

SYSTEM CAPABILITIES MATRIX

- ' **Multi-Tenancy** Relational database multi-tenancy with strict foreign key cascading and context scoping.
- ' **RAG Pipeline** Ingests PDF, DOCX, XLSX, TXT into 768-dim text-embedding-004 vectors with Cosine Reranking.
- ' **Streaming Engine** Direct Server-Sent Events (SSE) stream via Google Gemini 2.0 Flash & 1.5 Pro (TTFT < 380ms).
- ' **Isolated Widget** Lightweight (<25KB) Shadow DOM chat widget with zero CSS bleeding and dynamic agent switcher.
- ' **Security Vault** Passwordless Nodemailer OTP login, AES-256-GCM secret vault, timing-safe hash comparison.
- ' **CR Ticketing** Threaded customer relations support system with priority matrix and Super Admin notes.

Author / Engineering Team:


BrainPlug Architecture & Core Infrastructure Group

Framework: Next.js 16 • Database: PostgreSQL 16 • AI: Google Gemini • Date: September 2026

Executive Summary & System Topology

BrainPlug is an enterprise-grade, multi-tenant conversational AI agent platform engineered to empower modern organizations to create, customize, ground, and deploy secure AI assistants in seconds. By coupling Google Gemini's advanced multimodal reasoning capabilities with an in-database vector retrieval-augmented generation (RAG) engine, BrainPlug delivers high-fidelity, hallucination-free customer support and operational automation.

The BrainPlug system is structured across four decoupled architectural layers, ensuring maximum scalability, zero host page style collisions, and robust multi-tenant data boundaries.



By avoiding external vector SaaS dependencies (e.g. Pinecone) and hosting dense vector embeddings directly inside PostgreSQL with Prisma transactions, BrainPlug achieves atomic data isolation, zero external API hops for vector retrieval, and instant transactional rollbacks.

SECTION 2

Multi-Tenancy & Data Isolation Model

2.1 Logical Multi-Tenancy Architecture

BrainPlug implements logical database multi-tenancy enforced through non-nullable `tenantId` foreign key discriminators across all domain tables. This architecture balances high operational density on shared compute with ironclad tenant data partitioning.

- **Foreign Key Cascading:** Deleting or modifying a tenant automatically cascades across its agents, documents, chunks, API keys, conversations, and audit logs.
- **Dynamic Context Resolution:** Every API request resolves tenant identity via `requireTenantAccess()` in `server/auth/context.ts`. Client users are strictly constrained to their tenant boundary.
- **Super Admin Clearance:** Users with the `SUPER_ADMIN` system role bypass tenant-specific `WHERE` clauses to inspect platform-wide metrics, manage models, and resolve cross-tenant support tickets.

2.2 Multi-Tenant Role-Based Access Control (RBAC)

System Role	Workspace Scope	Allowed Capabilities	Target User Type
SUPER_ADMIN	Global Platform	Inspect all tenants, configure Gemini models, view platform telemetry, manage system roles.	Platform Owner / DevOps
CLIENT_ADMIN	Tenant Workspace	Create/edit AI agents, upload knowledge docs, generate API keys, manage support tickets, view	Enterprise Customer Admin
CLIENT_USER	Tenant Workspace	Interact with AI agents, review conversation transcripts, view knowledge base documents.	Support Agent / Employee

2.3 Context Scoping Code Enforcement

```
// server/auth/context.ts - Dynamic Tenant Resolution
export async function requireTenantAccess(targetTenantId?: string) {
  const user = await requireAuth();
  if (user.role === 'SUPER_ADMIN') {
    return { user, tenantId: targetTenantId || user.tenantId || '' };
  }
  if (!user.tenantId || (targetTenantId && targetTenantId !== user.tenantId)) {
    throw new AppError('Tenant access denied', 'TENANT_ACCESS_DENIED', 403);
  }
  return { user, tenantId: user.tenantId };
}
```

SECTION 3

RAG Knowledge Ingestion & Vector Search

3.1 Document Ingestion & Semantic Chunking

To provide verifiable grounding for conversational AI agents, BrainPlug includes a high-throughput document ingestion engine capable of parsing diverse enterprise formats and generating semantic vector embeddings.

- **1. Multi-Format Text Extraction:** Extracts structured text from PDF (pdf-parse), DOCX (mammoth), XLSX spreadsheets (xlsx), and TXT/Markdown.
- **2. 500-Token Sliding Window:** Splits text into 500-token chunks with 50-token overlap, preserving context continuity across boundaries.
- **3. text-embedding-004 Vectorization:** Generates 768-dimensional dense float32 vector arrays using Google's state-of-the-art embedding model.
- **4. Relational Persistence:** Stores chunks with embedding JSON arrays in document_chunks table with tenantId & agentId composite indexes.

3.2 Real-Time Cosine Similarity Search Math

At inference time, the user's incoming message is converted to a 768-dimensional query vector. BrainPlug computes the Cosine Similarity against all active chunks assigned to the targeted agent:

```
// Cosine Similarity Formula:
similarity(Q, C) = (Q · C) / (||Q|| * ||C||) = " (Q_i * C_i) / (" " (Q_i^2) * " " (C_i^2))

// Selection Criteria:
1. Score Threshold: Minimum similarity score >= 0.40 (filters irrelevant noise)
2. Top-K Selection: Top 5 highest-ranking passages injected into Gemini System Prompt
3. Latency: Average vector search execution in under 18ms across 10,000 chunks
```

GROUNDING & ANTI-HALLUCINATION GUARDRAILS

When document chunks are retrieved, the agent's system prompt is dynamically rewritten with explicit instructions to prioritize source context and include citations. If no relevant chunks match the query, the agent gracefully acknowledges domain boundaries.

SECTION 4

Streaming Inference & Shadow DOM Widget

4.1 Server-Sent Events (SSE) Streaming Protocol

BrainPlug leverages direct Server-Sent Events (SSE) streaming from Google Gemini 2.0 Flash and 1.5 Pro. This achieves a Time-to-First-Token (TTFT) under 380ms, delivering immediate visual responsiveness.

Event Type	Payload Structure	Client Processing Behavior
start	{ "type": "start", "conversationId": "...", "sources": [...] }	Initializes conversation container, renders citation source badges.
token	{ "type": "token", "content": "text_chunk" }	Appends raw markdown token to active stream, auto-scrolls chat window.
done	{ "type": "done", "latencyMs": 342, "usage": {...} }	Finalizes message bubble, commits usage telemetry, enables input field.
error	{ "type": "error", "message": "Quota exceeded" }	Renders user-friendly alert bubble, offers support ticketing link.

4.2 Embeddable Shadow DOM Chat Widget (widget.js)

The client chat widget is distributed as an ultra-lightweight (<25KB), standalone vanilla JavaScript file (`widget.js`) with zero third-party framework dependencies. It encapsulates the full chat UI within a closed Shadow DOM container (`mode: 'closed'`), preventing any host website CSS leakage or style collision.

```
<!-- 1-Line Embed Code for Any Website -->
<script
  src="https://api.brainplug.ai/widget.js"
  data-agent-id="00000000-0000-0000-0000-000000000001"
  data-api-key="bp_live_a1b2c3d4e5f6g7h8i9j0..."
  data-primary-color="#7c3aed"
  data-position="bottom-right"
  defer>
</script>
```

- **Dynamic Agent Switcher:** The client workspace portal includes an instant agent toggle allowing administrators to preview and live-test multiple configured AI bots directly within the dashboard.
- **WYSIWYG Live Customizer:** Real-time control over brand colors, typography, launcher icons (Message, Bot, Sparkles), corner radius (0-24px), and mobile bottom-sheet transitions.

SECTION 5

Security Architecture & Support Ticketing

5.1 Cryptographic Architecture & Authentication

Security is engineered at every layer of BrainPlug, encompassing passwordless authentication, secret encryption, timing-safe validation, and domain restriction guards.

- **Passwordless OTP Authentication:** Users register and log in via 6-digit single-use OTP codes dispatched through Nodemailer SMTP relays. OTPs are stored as SHA-256 hashes with 5-minute expiry.
- **Timing-Safe Hash Comparison:** All token and OTP verifications utilize `crypto.timingSafeEqual()` to eliminate side-channel timing attack vulnerabilities.
- **AES-256-GCM Secret Vault:** API secret keys and third-party credentials are encrypted using authenticated symmetric AES-256-GCM with randomized 16-byte IVs and 16-byte auth tags.
- **Allowed Domains Security Guard:** Chat widget endpoints validate the Origin header against the tenant's `allowed_domains` table, preventing unauthorized embedding on third-party sites.

5.2 Customer Relations (CR) & Support Ticketing

BrainPlug features an integrated multi-tier support ticketing system facilitating direct communication between client organizations and platform administrators.

Priority Level	Target SLA	Category Classification	Escalation Behavior
URGENT	1 Hour	TECHNICAL_ISSUE, BILLING	Instant email alert to on-call Super Admins.
HIGH	6 Hours	AGENT_CONFIG, RAG_KNOWLEDGE	High-priority queue placement in admin portal.
MEDIUM	24 Hours	GENERAL, FEATURE_REQUEST	Standard operational review queue.
LOW	48 Hours	GENERAL	Routine non-blocking inquiry workflow.

SUPER ADMIN INTERNAL NOTES

The ticketing module includes a confidential Internal Note toggle (`isInternalNote: true`) enabling administrators to collaborate and document investigation details without exposing internal comments to client users.

SECTION 6

Database Schema & Developer API

6.1 Core PostgreSQL Relational Entity Schema

Entity Model	Primary Keys & Discriminators	Relations & Cascading	Core Attributes
tenants	id (UUID), slug (Unique)	agents, documents, users, tickets	name, company_name, status, logo_url
users	id (UUID), email (Unique)	tenantRoles, sessions, tickets	full_name, mobile, status, email_verified
agents	id (UUID), tenant_id (FK)	widgetConfig, documents, apiKeys	name, system_prompt, type, status
documents	id (UUID), tenant_id, agent_id	chunks (Cascade Delete)	file_name, file_size, mime_type, status
document_chunks	id (UUID), document_id (FK)	document (Cascade Delete)	content, embedding (JSON), token_count
api_keys	id (UUID), key_hash (Unique)	tenant, agent, creator	name, key_prefix, status, scopes, revoked_at
tickets	id (UUID), ticket_number (Unique)	messages (Cascade Delete)	title, description, category, priority, status

6.2 Developer Streaming Chat API Specification

```
POST /api/v1/chat – Real-Time Inference Endpoint
Headers: Content-Type: application/json | Authorization: Bearer bp_live_...

// Request Body:
{
  "agentId": "000000000-0000-0000-0000-0000000000001",
  "message": "What are the return conditions?",
  "stream": true
}

// Response (Server-Sent Events text/event-stream):
data: {"type":"start","conversationId":"3fa85f64...", "sources":[{"fileName":"policy.pdf"}]}
data: {"type":"token","content":"You can return items within 30 days..."}
data: {"type":"done","latencyMs":342, "usage":{"totalTokens":184}}
```

SECTION 7

Performance Benchmarks & Strategic Roadmap

7.1 Measured Performance Benchmarks

Pipeline Operation	BrainPlug Performance	Industry Standard	Architectural Advantage
Time to First Token (TTFT)	340 ms	950 ms	Zero intermediate buffer, SSE direct piping
Vector Cosine Search	18 ms	85 ms	In-database PostgreSQL JSON array vector math
Document Ingestion Rate	250 pages / min	60 pages / min	Parallel token chunker with worker streams
Widget Bundle Footprint	< 25 KB	180 KB	Pure vanilla JavaScript + Shadow DOM isolation
API Auth & Scoping Overhead	< 4 ms	25 ms	In-memory timingSafeEqual & SHA-256 caching

7.2 Edge-Case Resilience & Hardening Matrix

- **Empty Tenant Scoping:** Super Admin route handlers gracefully evaluate cross-tenant lookups without throwing false 404 Not Found errors.
- **Token Overflow Protection:** Automatic prompt sliding-window truncation prevents inputs from exceeding Gemini maximum context boundaries.
- **Corrupt File Quarantine:** Non-blocking error handling logs parsing failures while preserving healthy documents in the batch.

7.3 Strategic Future Roadmap

- **Q3 2026: Multimodal Vision RAG:** Ingestion and semantic search across technical blueprints, schematics, and product catalogs.
- **Q4 2026: Real-Time Voice Agents:** Low-latency WebRTC bidirectional audio streaming powered by Gemini Multimodal Live API.
- **Q1 2027: Autonomous Multi-Agent Swarms:** Collaborative teams of specialized agents (Sales, Support, Technical, Billing) routing customer queries dynamically.
- **Q2 2027: Enterprise Webhook Engine:** Real-time webhook dispatchers for HubSpot, Salesforce, Slack, Zendesk, and ServiceNow integrations.

